

At batch 128, the median counter reading is at least 3× the MFU on every GPU tested

Annex A.2.1 is linear in U, so each ratio is the factor by which reported training compute changes with the reading of "GPU utilization". Medians over models, bf16 and fp16; 11-token sequences.

