


On an H100 at 2k–8k tokens, the counter reads 2.4–3.1× the MFU

Same training steps. The counter stays at 99.8–100% while MFU moves between 32% and 41%.
Annex A.2.1 is linear in U, so each ratio is the factor by which reported compute changes.



Measured 16 Sep 2026 on one NVIDIA H100 PCIe (Scaleway H100-1-80G). TinyLlama-1.1B architecture, random tokens, bf16, torch.compile. Peak 756.45 TFLOP/s dense BF16.